

Homework 5

Problem 1 - *Deep Reinforcement Learning* 25 points

This question is based on Deep RL concepts discussed in Lecture 8. You need to refer to the papers by Mnih et al., Nair et al., and Horgan et al. to answer this question. All papers are linked below.

1. Explain the difference between episodic and continuous tasks? Given an example of each. (2)
2. What do the terms exploration and exploitation mean in RL ? Why do the actors employ ϵ -greedy policy for selecting actions at each step? Should ϵ remain fixed or follow a schedule during Deep RL training ? How does the value of ϵ help balance exploration and exploitation during training. (2+2+1+1)
3. How is the Deep Q-Learning algorithm different from Q-learning ? You will follow the steps of Deep Q-Learning algorithm in Mnih et al. (2013) page 5, and explain each step in your own words. (3)
4. What is the benefit of having a target Q-network ? (3)
5. How does experience replay help in efficient Q-learning ? (3)
6. What is prioritized experience replay ? (2)
7. Compare and contrast GORILA (General Reinforcement Learning Architecture) and Ape-X architecture. Provide 3 main similarities and differences. (6)

References

- Mnih et al. Playing Atari with Deep Reinforcement Learning. 2013
Available at <https://arxiv.org/pdf/1312.5602.pdf>
- Nair et al. Massively Parallel Methods for Deep Reinforcement Learning. 2015
Available at <https://arxiv.org/pdf/1507.04296.pdf>
- Horgan et al. Distributed Prioritized Experience Replay. 2018
Available at <https://arxiv.org/pdf/1803.00933.pdf>

Problem 2 - TTA metric, Stability, Generalization 25 points

This question is based on efforts around DL benchmarking and standardization of performance metrics. We will study TTA metric properties using CIFAR10 dataset. Observe that, MLPerf and DAWNBench both use Imagenet1K for benchmarking but we will use CIFAR10, purely for convenience. In particular, we want to study (i) TTA stability, and (ii) Generalization performance of models optimized for TTA. A similar study was done in Coleman et al. using MLPerf and DAWNBench models. You should read this paper (especially Section 4.1 and 4.2) to have a better understanding of what you will be doing in this question. To study stability we will compute the coefficient of variation of TTA on different hardware by running several times the same training job. You will train Resnet-50 model in Pytorch with 2 different hardware types (V100, TPU pod) using CIFAR10. You will use a batch size of 128 and run each job for 350 epochs. From the training logs get the total (wall-clock) time to reach 92% accuracy for each of the runs. **We need data from at least 5 different training runs for the same configuration. This data collection can be done in a group of 5 students. Each of the student (in a group of 5) can run the same two training jobs, one with V100 and one with a TPU pod. If you decide to collaborate in the data collection please clearly mention the name of students involved in your submission.**

1. Calculate the coefficient of variation of TTA for both the hardware configurations. Compare the value you obtain with that reported in Table 3(a) in the paper by Coleman et al for Resnet-50, 1xTPU. (5)

Homework 5

2. Collect 5 images from the wild for each of the 10 categories in CIFAR10 and manually label them. **Again in a group of 5 students, each one of you can collect 5 images for 2 categories. In this way each one of you need to collect only 10 images (5 for each of the 2 categories that you choose).** These images should not be from CIFAR10 dataset. Next test the trained models for these 50 images. What is the accuracy obtained from each of the 10 trained models ? Quantify the mean and standard deviation of accuracy obtained using the 10 models for TPU and the 10 models for V100. (5+5)
3. [BONUS] Measure the GPU utilization using `nvidia-smi` while training with V100 and report the average GPU utilization over a period of 3 mins of training. Is the GPU utilization close to 100% ? (5)
4. [BONUS] If the GPU utilization is low, what can you do to increase the GPU utilization? Try your trick(s) and report if you are successful or not in driving GPU utilization close to 100%. (5)

References

- Coleman et al. Analysis of DAWNBench, a Time-to-Accuracy Machine Learning Performance Benchmark. Available at <https://cs.stanford.edu/matei/papers/2019/sigops-osr-dawnbench-analysis.pdf>

Problem 3 - SSD, ONNX model, Visualization, Inferencing 40 points

In this problem we will be inferencing SSD ONNX model using ONNX Runtime Server. You will follow the github repo and ONNX tutorials (links provided below). You will start with a pretrained Pytorch SSD model and retrain it for your target categories. Then you will convert this Pytorch model to ONNX and deploy it on ONNX runtime server for inferencing.

1. Download pretrained pytorch MobilenetV1 SSD and test it locally using Pascal VOC 2007 dataset. Show the test accuracy for the 20 classes. (5)
2. Select any two related categories from Google Open Images dataset and finetune the pretrained SSD model. Examples include, Aircraft and Aeroplane, Handgun and Shotgun. You can use `open_images_downloader.py` script provided at the github to download the data. For finetuning you can use the same parameters as in the tutorial below. Compute the accuracy of the test data for these categories before and after finetuning. (5+5)
3. Convert the Pytorch model to ONNX format and save it. (5)
4. Visualize the model using net drawer tool. Compile the model using `embed_docstring` flag and show the visualization output. Also show doc string (stack trace for PyTorch) for different types of nodes. (10)
5. Deploy the ONNX model on ONNX runtime (ORT) server. You need to set up the environment following steps listed in the tutorial. Then you need make HTTP request to the ORT server. Test the inferencing set-up using 1 image from each of the two selected categories. (5)
6. Parse the response message from the ORT server and annotate the two images. Show inferencing output (bounding boxes with labels) for the two images. (5)

For part 1, 2, and 3, refer to the steps in the github repo. For part 4 refer to ONNX tutorial on visualizing and for 5 and 6 refer to ONNX tutorial on inferencing.

References

Homework 5

- Github repo. Shot MultiBox Detector Implementation in Pytorch.
Available at <https://github.com/qfgaohao/pytorch-ssd>
- ONNX tutorial. Visualizing an ONNX Model.
Available at <https://github.com/onnx/tutorials/blob/master/tutorials/VisualizingAModel.md>
- ONNX tutorial. Inferencing SSD ONNX model using ONNX Runtime Server.
Available at <https://github.com/onnx/tutorials/blob/master/tutorials/OnnxRuntimeServerSSDModel.ipynb>
- Google. Open Images Dataset V5 + Extensions.
Available at <https://storage.googleapis.com/openimages/web/index.html>
- The PASCAL Visual Object Classes Challenge 2007.
Available at <http://host.robots.ox.ac.uk/pascal/VOC/voc2007/>

Problem 4 - ML Cloud Platforms 30 points

In this question you will analyze different ML cloud platforms and compare their service offerings. In particular, you will consider ML cloud offerings from IBM, Google, Microsoft, and Amazon and compare them on the basis of following criteria:

1. Frameworks: DL framework(s) supported and their version. (4)
2. Compute units: type(s) of compute units offered. (2)
3. Storage: storage type(s) supported. (3)
4. Reinforcement Learning: explicit support for RL training offered ? If yes, provide the url and service details. (5)
5. Model lifecycle management: tools supported to manage ML model lifecycle. (2)
6. Monitoring: availability of resource usage monitoring data to the user. (2)
7. Visualization during training: application logs, metrics like accuracy and throughput (2)
8. Elastic Scaling: support for elastic scaling compute resources of an ongoing job. (2)
9. Training job description: training job description file format. Show how the same training job is specified in different ML platforms. Identify similar fields in the training job file for the 4 ML platforms through an example. (8)

Problem 5 - *Transfer Learning using Kubeflow on Google Cloud* 40 points

In this problem we will follow Kubeflow codelab (link below). You will follow the steps as outlined in the codelab. The only difference is that instead of training MNIST model you will start with a pretrained Resnet-50 model on Imagenet1K and then finetune it for the dataset you chose from visual-decathlon for question1 of Hw4. Note here the training will be done on the cluster using Kubeflow. For finetuning use the learning rate that gave you the best accuracy in question1 of Hw4. For each step below you need to show the commands and output (if any).

1. *Getting Started*: Follow steps to set environment variables, install Kustomize, enable the API, and set up a Kubeflow cluster. (6)

Homework 5

2. *Training*: Follow steps to set up a storage bucket, build the container image using pretrained model and python code for finetuning, test locally, train (finetune) on the cluster. Before building the container image you need to use Kustomize to change the parameters in the manifest file like batch size, number of epochs, learning rate, and other hyperparameters you will use. Show the manifest file that you will use for training. At the end of training on the cluster check the content of your storage bucket. Show a screen shot of the contents as in the codelab. (4+4+2+4+4)
3. *Serving*: Use Kustomize to edit the manifest file for serving and point to the GCS bucket containing the trained model. Then deploy the server to the cluster. (4)
4. *Deploying the UI*: Follow steps to create a web interface to interact with the deployed model. Create a script to (similar to `mnist_client.py`) to interact with the Tensorflow server through gRPC. Next deploy the web UI, load the web UI in your browser and show the web UI screenshot. (4+4)
5. *Inferencing*: Send two images from the test set to the service (using the client script) and get the top-5 predicted classes. Show the service response. (4)

References

- Codelab. Introduction to KubeFlow on Google Kubernetes Engine.
Available at <https://codelabs.developers.google.com/codelabs/kubeflow-introduction/>